

Bayes' rule

Giorgio Corani - (IDSIA, SUPSI)

Bayesian Data analysis and Probabilistic
Programming

- Chap. 2 of *Bayes Rules! An Introduction to Applied Bayesian Modeling*
 - <https://www.bayesrulesbook.com/chapter-2.html>

- An event is a set of outcomes of an experiment to which a probability is assigned.
- Examples of events and probabilities related to a fair dice:
 - E_1 : the dice lands 4 ($P=1/6$)
 - E_2 : the dice lands odd: 1 or 3 or 5. ($P=1/2$)
 - E_3 : the dice lands 2 or 4. ($P=1/3$)

Conditional vs unconditional probability

- Let A and B be two events.
- $P(A)$: unconditional probability of A . It measures the probability of observing A , without any knowledge of B .
- $P(A|B)$: conditional probability of A given B : probability of observing A once B occurred (probability of A given B).

- A is not observed.
- Comparing $P(A|B)$ vs $P(A)$ reveals how the observation of B informs us about A .
- $P(A|B)$ can be larger, smaller or equal to $P(A)$.

- Probability of joining an orchestra, given that one practices clarinet every day:

$$P(\text{orchestra} \mid \text{practice}) > P(\text{orchestra})$$

- Probability of getting the flu given that one washes thoroughly his hands:

$$P(\text{flu} \mid \text{wash hands}) < P(\text{flu})$$

$$P(A|B) \neq P(B|A)$$

- Roughly 100% of puppies are adorable:

$$P(\text{adorable} \mid \text{puppy}) = 1$$

- But an adorable object is not necessarily a puppy:

$$P(\text{puppy} \mid \text{adorable}) < 1$$

$P(B|A) > P(B)$ **or** $P(B|A) < P(B)$?

- B = you just finished reading a book from a certain author and you enjoyed it.
- A = you will also enjoy the newest novel from the same author.

$P(B|A) > P(B)$ **or** $P(B|A) < P(B)$?

- B = it's 0 degrees Celsius tonight.
- A = tomorrow it will be very warm.

$$P(B|A) > P(B) \text{ or } P(B|A) < P(B)?$$

Consider a woman who is mother of two children.

- B = the first child is a boy
- A = the second child will be a girl

Independent events

- A and B are independent if the occurrence of B doesn't tell us anything about the occurrence of A :


$$P(A|B) = P(A)$$

- Marginal and conditional probability are identical.
- For instance:
 - event A : rider Q wins the next motoGP race
 - event B : my coin lands tails
 - Since the coin does not affect the performance of rider Q,
 $P(A|B) = P(A)$.

Motivating example

- According to a study, 40% of the articles taken from a certain website are fake news and 60% are real.
- The usage of an exclamation point in the title is however uncommon in real articles.
 - 26.67% of fake news titles use an exclamation point
 - only 2.22% of real news titles use it.

- Given that an article contains the exclamation point, which is its probability of being fake news?



Notation

- We denote random variable with uppercase letters and their specific values by lowercase letters.
- Let us consider discrete variables, which have a finite set of possible outcomes.
- $P(Y)$: probability distribution for variable Y , i.e, a table containing as many elements as the possible values of Y .
 - If Y is discrete, the distribution is also called a probability mass function (pmf).

- Specific values are denote by lowercase; for instance, y .
- $p(y)$ denotes the probability of $P(Y = y)$, where $\sum_y p(y) = 1$.
- We use moreover the shortcut notation
$$P(x|y) = P(X = x \mid Y = y)$$

- From now we will deal with random variables, not event.
- Consider the random variable A (article) with possible states $\{real, fake\}$.
- Prior probability, i.e., before having seen the article:
 - $P(A = fake) = 0.4$
 - $P(A = real) = 0.6$

A further random variable

- Variable B refers to whether the article title contains or not an explanation point.
- Its possible states are $\{!, \sim!\}$

Conditional probability

- The exclamation point in the title is more compatible with fake news than with real news.
- If an article is fake, there is 26.67% chance it uses exclamation points in the title.
- The conditional probability $P(!|fake)$ is 0.2667.
- The conditional probability $P(!|real)$ is instead 0.0222.

- The likelihood measures the compatibility of the observation (article contains exclamation point) with the different values of the unobserved variable (article is fake or real).
- $P(A|b)$: table containing the probability of the each value of A , given the observation $B = b$.

Prior probabilities and likelihoods of fake news.

	fake	real
prior probability	$P(\text{fake})=0.40$	$P(\text{real})=0.6$
likelihood	$P(! \mid \text{fake})=0.26$	$P(! \mid \text{real})=0.02$

- Prior probabilities add up to 1 but the likelihoods do not.
- The likelihood function is not a probability function; instead it is a way to measure the compatibility of the observation (title contains !) with the different states of the unobserved variable A (fake or real).

	fake	real
!		
~!		

- The joint probability $P(a, b)$ is the probability of observing both $A = a$ and $B = b$.
- It is defined as

$$P(a, b) = P(a|b)P(b)$$

$$\begin{aligned}P(\text{fake}, !) &= P(\text{fake})P(! \mid \text{fake}) \\&= 0.4 \cdot 0.2667 = 0.1067\end{aligned}$$

	fake	real
!	0.1067	
~!		

$$\begin{aligned}P(\text{fake}, \sim !) &= P(\text{fake})P(\sim ! \mid \text{fake}) \\&= 0.4 \cdot (1 - 0.2667) = 0.2933\end{aligned}$$

	fake	real
!	0.1067	
$\sim !$	0.2933	

$$\begin{aligned}P(\text{real}, !) &= P(\text{real})P(! \mid \text{fake}) \\ &= 0.6 \cdot 0.0222 = 0.0133\end{aligned}$$

	fake	real
!	0.1067	0.0133
~ !	0.2933	

$$\begin{aligned}P(\text{real, no !}) &= P(\text{real})P(\text{no !} \mid \text{fake}) \\&= 0.6 \cdot (1 - 0.0222) \\&= 0.6 \cdot 0.9778 = 0.5867\end{aligned}$$

	fake	real
!	0.1067	0.0133
~ !	0.2933	0.5867

- The joint probability sums up to 1.

Marginal distribution and marginal probability

- The marginal distribution of A is obtained by summing the joint distribution over all states of B .

	fake	real
!	0.1067	0.0133
$\sim !$	0.2933	0.5867
marginal	0.4	0.6

- The marginal distribution of A is {fake=0.4; real=0.6}.
- Moreover,

$$P(A = \text{fake}) = P(\text{fake}, !) + P(\text{fake}, \sim !)$$

.

Marginal distribution and marginal probability

- The marginal distribution of B is obtained by summing the joint distribution over all states of A .

	fake	real
!	0.1067	0.0133
~ !	0.2933	0.5867

- The marginal distribution of B is $\{! = 0.12; \sim ! = 0.88\}$.
- A randomly chosen article has 0.12 probability of containing the exclamation point in the title.

Sum rule

- Given the joint distribution $P(X, Y)$
- The marginal probability $P(X = x)$ is given by the **sum rule**

$$P(x) = \sum_y P(x, y)$$

- The summation is over all possible values of Y .

Computing marginal probabilities

- Assume you want to compute $P(b)$ from $P(A)$ and $P(B|A)$, rather than from $P(A, B)$.
- The marginal $P(B = b)$ is $P(b) = \sum_a P(b, a) = \sum_a P(b | a)P(a)$.

$$\begin{aligned}P(!) &= P(\text{real}, !) + P(\text{fake}, !) \\&= P(\text{real})P(! | \text{real}) + P(\text{fake})P(! | \text{fake}) \\&= 0.4 \cdot 0.0222 + 0.6 \cdot 0.2667 \\&= 0.1067 + 0.1600 = 0.2667\end{aligned}$$

Posterior probability that the article is fake

- We want the posterior probability of the article being fake given that it uses exclamation points.

$$P(\text{fake} | !) = \frac{P(\text{fake})P(! | \text{fake})}{P(!)} = \frac{0.4 \cdot 0.2667}{0.12} = 0.889$$

Bayes' theorem - recap

- Unobserved variable A , whose set of possible states is denoted by \mathcal{A} .
- Observable variable B

Requires:

- prior distribution $P(A)$
- conditional distribution $P(B|A)$
- Given $B = b$, the posterior distribution of A is:

$$P(A|b) = \frac{P(A)p(A | b)}{P(b)} = \frac{P(A)p(A | b)}{\sum_{a \in \mathcal{A}} P(a, b)}$$
$$P(b) > 0$$

- $P(A)$ represents how likely are the different values of A , according to our beliefs, before we see any data.

- $P(A \mid b)$ represents the distribution of the possible outcome A given $B = b$.
- When we have a specific observation b , we get the likelihood $P(a|b)$ for each $a \in \mathcal{A}$.
- The likelihood is not a probability distribution; it does not sum to one.

Marginal likelihood

- The denominator of Bayes' rule is a normalizing constant, referred to as the marginal likelihood.
- It marginalizes the likelihood over the states of the unobserved variable A :

$$P(b) = \sum_a p(a)P(b \mid a) = \sum_a P(a, b)$$

- You decide to take a diagnostic test to check if you have contracted Covid. You want to make inference about your health H whose possible states are:
 - infected
 - healthy

The diagnostic test

The test T can be either:

- positive
- negative

We want to determine the probability distribution $P(H|positive)$.

Test performance: conditional probability

Assume the conditional probability of the test outcome, given an infected person, to be:

	test negative	test positive
$P(T \mid \text{infected})$	0.125	.875

- Thus the probability of a positive test for an infected persons is 87.5%.

Test performance: conditional probability

Assume the conditional probability of the test outcome, given an healthy person, to be:

	test negative	test positive
$P(T \mid \text{healthy})$.975	.025

- Thus the probability of a negative test for a healthy patient, 97.5%

Prior

- The *prevalence* is the percentage of persons affected by the disease.
- The covid prevalence in New York City 2020 was 10%.

	$H=\text{healthy}$	$H=\text{infected}$
probability	0.9	.1

- This is our prior, before observing the outcome of the test.

$P(\text{infected} \mid \text{positive})$

$$P(\text{infected} \mid \text{positive}) = \frac{\overbrace{P(\text{infected})}^{\text{prior: prevalence}} \overbrace{P(\text{positive} \mid \text{infected})}^{\text{likelihood}}}{\underbrace{P(\text{positive})}_{\text{prob of observing a positive test: marginal likelihood}}}$$

Denominator, a.k.a. marginal likelihood

- Total probability of having a positive test:
 - probability of testing positive while infected + ...
 - probability of testing positive while healthy

$$\begin{aligned}P(\text{positive}) &= P(\text{positive, infected}) + P(\text{positive, healthy}) \\&= P(\text{positive} \mid \text{infected})P(\text{infected}) + P(\text{positive} \mid \text{healthy})P(\text{healthy}) \\&= 0.875 \times 0.1 + 0.025 \times 0.9 \\&= 0.11\end{aligned}$$

$$\begin{aligned}P(\text{infected}|\text{positive}) &= \frac{P(\text{infected})P(\text{positive} \mid \text{infected})}{P(\text{positive})} \\&= \frac{0.1 \times 0.875}{0.11} \\&= 0.795\end{aligned}$$

- The posterior probability of being healthy is:

$$P(\text{healthy} \mid \text{positive}) = 1 - .795 = .205$$

- The positive test increases your probability of being infected of about 8 times.

Exercise

- Work out the probability of being infected if you test negative
 - 0.014

Yet another exercise

- We have two coins:
 - the first coin lands heads or tails with equal probability
 - the second coin is rigged and always lands heads.
- We take one coin at random and we get heads. What is the probability that this coin is the rigged one?

Since the coins are randomly chosen, the prior is:

$$P(\text{fair}) = P(\text{rigged}) = 0.5$$

The likelihood is:

$$P(\text{head} \mid \text{fair}) = 0.5$$

$$P(\text{head} \mid \text{rigged}) = 1$$

The posterior probability of the coin being rigged is:

$$\begin{aligned}P(\text{rigged} \mid \text{head}) &= \frac{P(\text{rigged})P(\text{head} \mid \text{rigged})}{P(\text{head})} \\&= \frac{0.5 \times 1}{P(\text{head})}\end{aligned}$$

Computing the denominator

$$P(\text{head}) = P(\text{head, rigged}) + P(\text{head, fair})$$

$$\begin{aligned} P(\text{head, rigged}) &= P(\text{head}|\text{rigged})P(\text{rigged}) \\ &= 1 \times 0.5 \end{aligned}$$


$$\begin{aligned} P(\text{head, fair}) &= P(\text{head}|\text{fair})P(\text{fair}) \\ &= 0.5 \times 0.5 \end{aligned}$$

$$P(\text{head}) = 0.5 \times 1 + 0.5 \times 0.5 = 0.75$$

The posterior probability of the coin which has landed tail to be the rigged one is:

$$\begin{aligned}P(\text{rigged} \mid \text{head}) &= \frac{P(\text{rigged})P(\text{head} \mid \text{rigged})}{P(\text{head})} \\&= \frac{0.5 \times 1}{0.75} = \frac{2}{3}\end{aligned}$$

Your turn



- You're given the choice of three doors:
 - Behind one door is a car
 - Behind the others, goats.
- You pick door 1 and the host, who knows what's behind the doors, opens door 3, which has a goat.
- He then says "Do you want to pick door No. 2?" Should you switch your choice?

■ TBD