Marcov Chain Monte Carlo (MCMC)

Giorgio Corani - SUPSI

Bayesian Data analysis and Probabilistic Programming

Credits

- J. Krushke, “Doing Bayesian Data Analysis”, Chapter 7

Non-conjugate likelihood

- Generally the likelihood is **not** conjugate to the prior.
- Grid approximation does not scale.
- If we need to compute the joint distribution of 6 parameters, each represented by 1000 states, we have 1000^6 states, too much for any computer.

MCMC

- MCMC methods address this type of problems. For simplicity, we show a single-parameter problem.
- We denote the prior by $p(\theta)$ and the likelihood by $p(D \mid \theta)$, where D denotes the data.
- The method avoids the direct evaluation of the difficult integral in the denominator of Bayes' rule:

$$p(\theta \mid D) = \frac{p(\theta)p(D \mid \theta)}{\underbrace{p(D)}_{\text{marg lik}}} = \frac{p(\theta)p(D \mid \theta)}{\int p(\theta)p(D \mid \theta)d\theta}$$

- The method approximates the posterior of θ by returning many samples.

Approximating a distribution with a (large) sample

- By randomly sampling a subset of people from a population, we can estimate the underlying tendencies in the entire population.
- The larger the sample, the better the estimation.
- The population from which we want to sample is the **posterior distribution of** θ.

Approximating a distribution by a large random sample

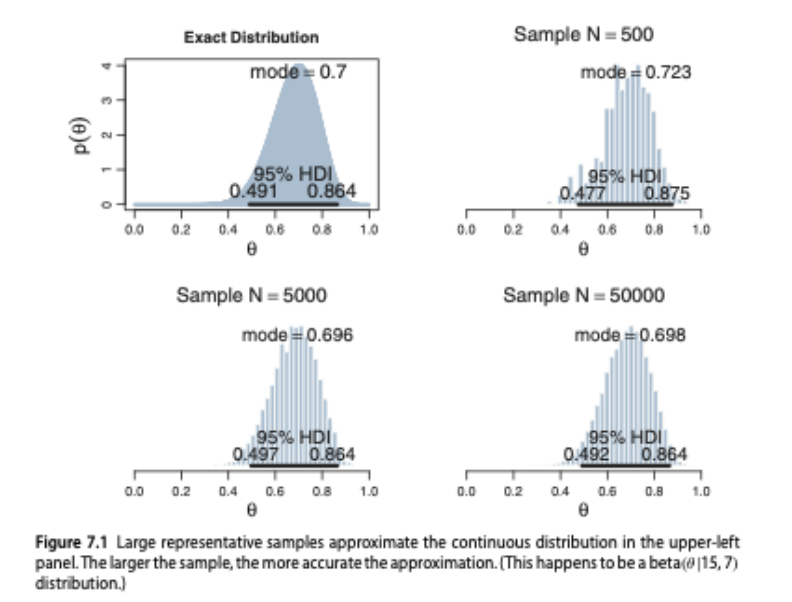

Figure 7.1 Large representative samples approximate the continuous distribution in the upper-left panel. The larger the sample, the more accurate the approximation. (This happens to be a beta(θ|15, 7) distribution.)

- Larger samples yield a more accurate histogram;
- The exact values (top-right) were obtained from the mathematical formula of the beta distribution.

The Metropolis algorithm: a simple example.

- We live in a chain of 7 islands.
- We have to do many travels, visiting each island proportionally to its population.
- We can move from an island to a neighboring one, or remain on the same island.
- The population of the different islands is 1000, 2000, 3000, ...,7000.

Making decisions

- Flip a coin to decide whether the *proposed* island is located east or west.
- Visit it **if** it has a larger population than the current one.
- Otherwise, visit it with probability $p_{\text{move}} = \frac{p_{\text{proposed}}}{p_{\text{current}}}$
- In the long run, each island is visited proportionally to its population!

Discussion

- At each time step, both the chosen direction and its acceptance are random.
- If the process were started over again, the specific trajectory would be different.
- Yet, in the long run the relative frequency of visits mimics in any case the target distribution.

Proposal

- We are at position θ_{current}.
- We randomly propose to move right (50%) or left (50%).
- The possible moves and the probability of proposing each is the *proposal distribution*.
- Our proposal distribution has only two values (left and right) with 50-50 probabilities.

Accepting the proposal

- If the target distribution is greater at the proposed position, we accept the proposed move.
  - we always move higher if we can.

- If the target distribution is less at the proposed position than at our current position, we accept the move with probability:

$$p_{\text{move}} = \frac{p_{\text{proposed}}}{p_{\text{current}}}$$

- We move to the proposed position with probability:

$$p_{\text{move}} = min\left(\frac{p(\theta_{\text{proposed}})}{p(\theta_{\text{current}})}, 1\right)$$

We use the *unnormalized* posterior

- The algorithm evaluates the ratio $\frac{p(\theta_{\text{proposed}})}{p(\theta_{\text{current}})}$, without computing the normalizing constant (i.e., the denominator) of Bayes rule.
- It does not require the absolute value of $p(\theta|D)$, but only the ratio between the density in different locations.
- We sample from $p(\theta)p(D|\theta)$ without normalizing by the (often) untractable marginal likelihood $p(D)$.
- In the example of islands-hopping, the target distribution was the unnormalized population of each island, not a normalized probability.

Discussion

- We want to sample from a target distribution.
- This is usually the unnormalized posterior distribution of θ, $p(\theta)p(D|\theta)$.
- Extensions for continuous values (see the following).
- Extenions for any number of dimensions (not covered).

Random walk

- The samples from the posterior are generated by a *random walk*
- The walk starts from a randomly chosen point where the distribution is non zero.
- At each time step we propose the move to a new position θ_{proposed}.
- We accept the move with probability

$$p_{\text{move}} = min\left(\frac{p(\theta_{\text{proposed}})}{p(\theta_{\text{current}})}, 1\right)$$

Metropolis algorithm applied to Bernoulli likelihood and beta prior

$$p(\theta \mid D) \propto p(D \mid \theta)p(\theta) = \theta^{a+y}(1-\theta)^{b+n-y}$$

- θ is a continuous parameter
- Proposal distribution: $\Delta\theta \sim N(0, \sigma)$.
- θ_{proposed} lies in an interval of $\pm 3\sigma$ around θ_{current}.
- Hence σ controls how far θ_{proposed} can be from θ_{current}.

Sampling Bernoulli likelihood and beta prior

Start from θ_0.

At each iteration:

- Draw $\Delta\theta \sim N(0, \sigma)$
- $\theta_{\text{proposed}} = \theta_{\text{current}} + \Delta\theta$

Probability of the move

We move to θ_{proposed} with probability:

$$p = \min\left(1, \frac{P(\theta_{\text{proposed}} \mid D)}{P(\theta_{\text{current}} \mid D)}\right)$$
$$= \min\left(1, \frac{\theta_{proposed}^{a+y}(1-\theta_{proposed})^{b+n-y}}{\theta_{current}^{a+y}(1-\theta_{current})^{b+n-y}}\right)$$

Application

- Consider the prior $p(\theta) = Beta(1, 1)$
- The Bernoulli likelihood $\theta^{14}(1 - \theta^6)$ corresponding to 20 tosses and 14 tails.
- The three columns use three different σ in the proposal distribution.

Results

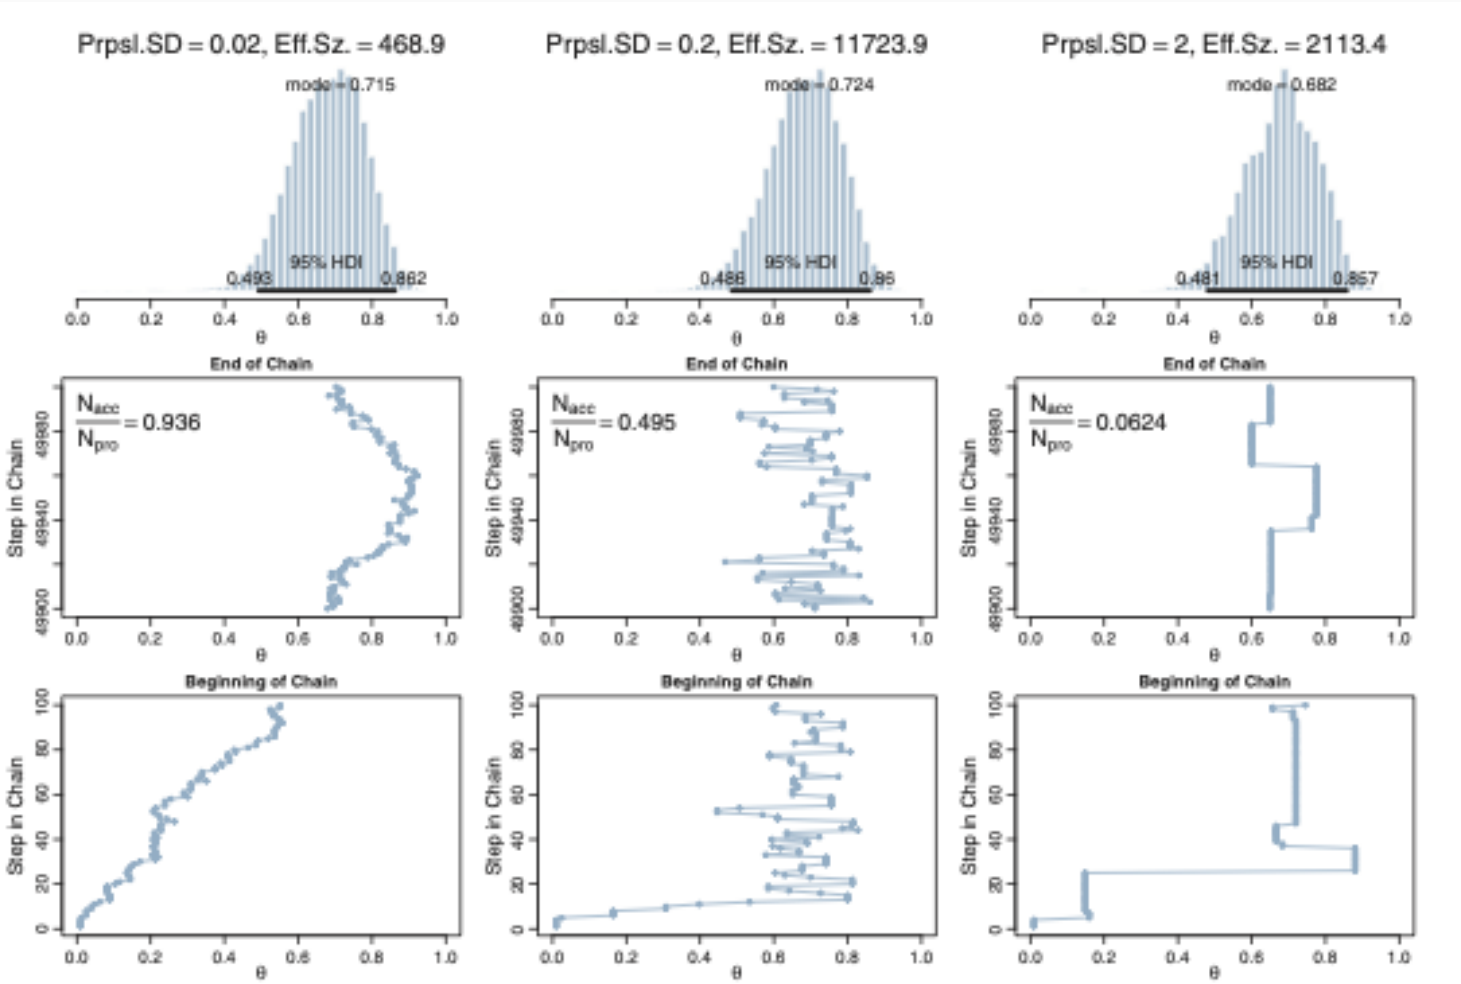

Figure 7.4 Metropolis algorithm applied to Bernoulli likelihood with beta($\theta|1, 1$) prior and $z = 14$ with $N = 20$. For each of the three columns, there are 50,000 steps in the chain, but for the left column, the proposal standard deviation (SD) is 0.02, for the middle column SD = 0.2, and for the right column SD = 2.0.

Left column (small σ)

- The left column uses a the small σ = 0.02.
- The successive steps in the chain make small moves
- The chain will require a very long chain to thorouhgly explore the posterior distribution.
- The effective size of this 50,000 step chain is only 468.9.

Right column (large σ)

- The proposed jumps are often far away from the bulk of the posterior distribution; the proposals are often rejected.
- The process accepts new values only occasionally, producing a very clumpy chain.
- In the long run, the chain will explore the posterior distribution thoroughly and produce a good representation, but it will require a very long chain.
- The effective size of this 50,000 step chain is only 2113.4.
- An acceptance ratio of about 0.5 usually provides the best effective sample size.
- Advanced implementations of the Metropolis algorithm automatically adjust the width of the proposal distribution.

Exercise 7.4 from Bayesrulebook - Tuning the Metropolis-Hastings)

rivedi

- Tune a Uniform proposal model with half-width w for a Metropolis-Hastings algorithm.
- Draw a trace plot for a tour where the Uniform proposal model uses a very small w
- Why is it problematic if w is too small, and hence defines the neighborhood around the current chain value too narrowly?

*Draw a trace plot for a tour where the Uniform proposal model uses a very large w. Why is it problematic if w is too large, and hence defines the neighborhood too widely?

- Draw a trace plot for a tour where the Uniform proposal model uses a w that is neither too small or too large.
- Describe how you would go about finding an appropriate half-width w for a Uniform proposal model.

Exercises

- Implement the island hopping algorithm? vedi codice in scripts
- Replicate the experiment about sampling the posterior beta.
- Replicate the experiment with a non-conjugate prior (triangular??).