

# The normal-normal model

Giorgio Corani

Bayesian Data Analysis and Probabilistic  
Programming

- Chap. 5 of *Bayes Rules! An Introduction to Applied Bayesian Modeling*
  - <https://www.bayesrulesbook.com/chapter-5.html>

- Let  $Y$  be a continuous random variable which can take values in  $(-\infty, \infty)$
- The variability of  $Y$  might be well represented by a Normal model

$$Y \sim N(\mu, \sigma^2)$$

# The Normal model

- The Normal pdf is

$$f(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y - \mu)^2}{2\sigma^2} \right]$$

- With:

$$E(Y) = \text{Mode}(Y) = \mu$$

$$\text{Var}(Y) = \sigma^2$$

$$\text{SD}(Y) = \sigma$$

## Standard deviation $\sigma$

- $\sigma$  provides a sense of scale for  $Y$ .
- Roughly 95% of  $Y$  values are within 2 standard deviations of  $\mu$ :

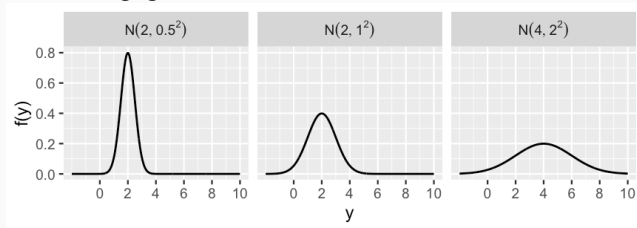
$$\mu \pm 2\sigma$$

- Roughly 99% of  $Y$  values are within 3 standard deviations of  $\mu$ :

$$\mu \pm 3\sigma$$

# The normal model

- The Normal model is bell-shaped and symmetric around  $\mu$ .
- As  $\sigma$  gets larger, the pdf becomes more spread out.
- Though a Normal variable is defined in  $(-\infty, \infty)$ , the plausibility of values that are more than 3 standard deviations  $\sigma$  from the mean  $\mu$  is negligible.



## Example

- The volume of the hippocampus (a part of the brain) is researched in studies about the effect of concussions.
- In the general population, both halves of the hippocampus have a volume between 3.0 and 3.5 cm<sup>3</sup>.
- Thus, the hippocampal volume is thought to vary, within the population, between 6 and 7 cm<sup>3</sup>.
- The average volume  $\mu$  is thought to be between 6.4 and 6.6 cm<sup>3</sup>.

## Normal prior

- Assuming symmetry, we formalize our prior information about  $\mu$  as:

$$\mu \sim N(\mu', \sigma_\mu)$$

which in this example yields :

$$\mu \sim N(6.5, 0.05)$$

- $\mu'$  is our prior guess on the value of  $\mu$ .
- $\sigma_\mu$  represents our uncertainty on the guess  $\mu'$ .
- According to this prior,  $\mu$  lies with 95% probability in (6.4, 6.6).
- We allow the volume to range over  $(-\infty, \infty)$ , but values beyond  $\mu \pm 3\sigma$  are given negligible probability.



- There is no single right prior, but different reasonable priors.

- We now define a model for the distribution of the observations.
- We make a *second* assumption of normality.
- The hippocampal volumes observed in  $n$  subjects  $(y_1, y_2, \dots, y_n)$  are normally distributed  $N(\mu, \sigma)$ .

- $\mu$  is the mean volume in the population.
- $\sigma$  expresses the spread of the measures within the population.
- We expect  $y$  to vary in (6-7); we interpret this interval as  $\mu \pm 2\sigma$ , hence it has length of  $4\sigma$ .
- We thus set  $\sigma=0.25$ .

- We moreover assume the observations  $y_1, \dots, y_n$  to be *independent* samples from  $N(\mu, \sigma)$ .
- This is realistic: the measure  $y_i$  tells us nothing about the measure  $y_{i+1}$  (assuming they refer to different subjects)

Assuming independence, the joint pdf of the  $n$  measures  $(y_1, y_2, \dots, y_n)$  is the product of the unique Normal pdfs  $f(y_i | \mu)$ :

$$f(\vec{y}|\mu) = \prod_{i=1}^n f(y_i|\mu) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y_i - \mu)^2}{2\sigma^2} \right].$$

- $\vec{y}$  is the vector containing the measures  $y_1, \dots, y_n$ .

# The Normal-Normal model

$$\mu \sim N(\mu', \sigma_\mu)$$

$$\vec{y} \sim N(\mu, \sigma)$$

- We treat  $\mu'$ ,  $\sigma_\mu$  and  $\sigma$  as fixed numbers.
- The only parameter of the model is  $\mu$ .
- Later we will treat also  $\sigma$  as a parameter.

## Your turn: normal likelihood functions

- For a Normal random sample  $y_i \sim N(\mu, \sigma)$  with  $\sigma=10$  we observe:

$$y_1, y_2, y_3 = (-4.3, 0.7, -19.4)$$

- Specify and plot the corresponding likelihood function of  $\mu$ .

## Conjugacy of the normal-normal model

- Denote the sample mean as  $\bar{y} = \frac{1}{n} \sum_i y_i$ .
- The posterior density of  $\mu$  is normal with updated parameters:

$$\mu | \vec{y} \sim N \left( \underbrace{\mu' \frac{\sigma^2}{n\sigma_\mu^2 + \sigma^2} + \bar{y} \frac{n\sigma_\mu^2}{n\sigma_\mu^2 + \sigma^2}}_{\text{posterior mean}}, \underbrace{\frac{\sigma_\mu^2 \sigma^2}{n\sigma_\mu^2 + \sigma^2}}_{\text{posterior variance}} \right).$$



## Posterior mean

$$\mu|\vec{y} \sim N\left(\mu' \frac{\sigma^2}{n\sigma_\mu^2 + \sigma^2} + \bar{y} \frac{n\sigma_\mu^2}{n\sigma_\mu^2 + \sigma^2}, \frac{\sigma_\mu^2 \sigma^2}{n\sigma_\mu^2 + \sigma^2}\right).$$

- The posterior mean is a weighted average of the prior mean  $\mu'$  and the sample mean  $\bar{y}$ .
- As  $n$  increases, the posterior mean converges to  $\bar{y}$ .
- As  $n$  increases, the posterior variance decreases.

## Your turn

- Which is the posterior mean, if we did 5 measures with  $\bar{y} = 6.7$ ?
- Which is the posterior mean, if we did 35 measures with  $\bar{y} = 6.7$ ?

## Your turn

- Let  $\mu$  be the average 3 p.m. temperature in Lugano.
- Your friend's prior understanding is that  $\mu$  is around 15 degrees Celsius, though might be anywhere between 5 and 25 degrees.
- To learn about  $\mu$ , he will analyze 1000 days of temperature data.
- Letting  $y_i$  denote the 3 p.m. temperature on day  $i$ , they'll assume that daily temperatures vary Normally around  $\mu$  with a standard deviation of 5 degrees.
- Formalize a normal-normal model.

## Treating $\sigma$ as a parameter


- A more sophisticated approach is to treat  $\sigma$  as a parameter.
- We assigning a prior to it; it should cover a wide range of plausible values for  $\sigma$ , leaving out however values that make no sense.
- In this case there is no closed-form expression of the posterior.

# Half-normal distribution

- $\sigma$  is strictly positive; a suitable prior is the *half-normal* distribution.
- The half-normal is a Gaussian restricted to positive values.
- Sample  $s$  from a half-normal are obtained by:
  - sampling from a normal distribution
  - applying the absolute value to the sampled values
  - $s \sim |N(0, \xi)|$ , where  $\xi$  is the standard deviation of the underlying normal. It is referred to as the *scale* of the half-normal.

# The half-normal distribution


- The HN pdf is asymmetric and right-skewed.
- It has long tails which are much larger than the median.



**Figure 1:** from wikipedia

## Effect of the scale parameter

- The half-normal pdf is characterized by a scale parameter (the standard deviation of the underlying normal).



## Tuning the half-normal distribution

- You can tune the scale of the HN by considering a plausible value of  $\sigma$ , and choose the scale so that it is close to the median of the HN.
- E.g., assume a plausible value of  $\sigma$  is 7.5.
- With 95% probability the measures are lie in an interval of  $\pm 15$  around the mean.
- But we are uncertain about this statement, as the interval could be well of  $\pm 30$ .



## Tuning the half-normal distribution

- We try different scales, until the median is about 7.5.
- Notice the long tails of the distribution, which allows to model to correct if our prior median guess (7.5) is underestimated.

```
pd.DataFrame(halfnorm.rvs(size=1000, scale=11)).describe()
```

```
##          0
## count  1000.000000
## mean    8.631128
## std     6.478797
## min     0.005291
## 25%     3.432332
## 50%     7.381743
## 75%    12.419468
## max    47.904591
```

## Probabilistic model with $\sigma$ as parameter

$\mu \sim N(\mu_\mu, \sigma_\mu)$       prior beliefs about  $\mu$

$\sigma \sim \text{Half-Normal}(\sigma_\sigma)$       prior beliefs about  $\sigma$

$y \sim N(\mu, \sigma)$       the observation are normally distributed  $\sigma$

- We cannot treat this model analytically, as the prior are no longer conjugates.
- We will implement it later via probabilistic programming.

- Try to define a probabilistic model of the distribution of height of adult males in Switzerland

- The mean height of the population could be 175.
- Keeping our prior broad, we state the mean height of the population to lie with 99% probability between 160 and 190 cm (the 99% interval roughly corresponds to  $\mu \pm 3\sigma$ ).
  - $\mu \sim N(175, 5)$

## Prior for $\sigma$

- We shall now assign a prior to  $\sigma$ . Within the population, we assume the height to lie with 99% probability between 100 and 250 (broad but realistic range).
- Hence the corresponding value of the standard deviation is  $(250-100)/6 = 25$ .

## Tuning the half-normal

- A half-normal distribution with scale 35 has roughly median 25:

```
pd.DataFrame(halfnorm.rvs(size=1000, scale=35)).describe()
```

```
##              0
## count  1000.000000
## mean    27.515002
## std     19.960007
## min      0.073953
## 25%     11.382362
## 50%     23.345175
## 75%     40.454872
## max     103.519415
```

## Likelihood (distribution of the data)

- Under the assumption of normality and independence, the likelihood is:

$$y \sim \mathcal{N}(\mu, \sigma)$$

- No further specification is required.

## The resulting model

$$\mu \sim N(175, 5)$$

$$\sigma \sim \text{half-normal}(35)$$

$$\vec{y} \sim N(\mu, \sigma)$$




## Solution for the exercise on normal likelihood functions

- Compute the likelihood as a function of  $\mu$  for a normal sample with  $\sigma=10$ , given the observations:

$$y_1, y_2, y_3 = (-4.3, 0.7, -19.4)$$

# Solution

```
#based on the observations, plausible values of mu  
#range between -30 and 20.  
mu = np.arange(-30, 20, 0.5)  
sigma = 10  
  
#a likelihood value for each value of mu  
lik = norm.pdf(-4.3, loc=mu, scale=sigma)  
  
#under independence, the likelihood of each observation multiplies  
lik = lik * norm.pdf(0.7, loc=mu, scale=sigma)  
lik = lik * norm.pdf(-19.4, loc=mu, scale=sigma)  
  
plt.figure(figsize=(10, 3))  
plt.plot(mu, lik)
```



- The values in the previous slide are numerically small. With more data, and more likelihood multiplication, it will become numerically untractable.
- For this reason it is numerically better to work with the log of the likelihood (log-likelihood).